

Recap

- What is an Image?
- What is a pixel?
- What is a Convolution?
- How do we store a continuous function in a computer?
 - Sampling
- What is aliasing?
 - Multiple continuous signals sharing the same sampling result
- What did Nyquist/Shannon do?
 - Theorem of sampling rate needed to represent a signal at a given frequency
- How does aliasing present itself in 2D graphics?

Review

- Properties of filters represented by convolution?
 - Linearly Separable
 - Shift Invariant
- Difference between moving average and convolution?
 - Convolution allows arbitrary normalization weights

Image Manipulation

Stephen J. Guy

Sep 18, 2014

Some Material Adopted from Peter Shirley and Steve Marshner & Greg Humphreys

Outline

- Warping
 - Scale, Rotate, Warps
- Sampling & Reconstruction
 - Resampling
 - Reconstruction Filters
- Image Manipulation
 - Filtering
 - Pixel Operation
- Quantization


Outline

- **Warping**
 - **Scale, Rotate, Warps**
- Sampling & Reconstruction
 - Resampling
 - Reconstruction Filters
- Image Manipulation
 - Filtering
 - Pixel Operation
- Quantization

5

Warping


- Moving Pixels of Images
 - Mapping
 - Resampling



6

Mapping


- Defines transformation between two images
 - Describe a destination (x,y) for every location (u,v) in the source
 - Vice-versa if it's invertable



7

Example Mapping - Scale


- Scale by factor:
 - $x = \text{factor} * u$
 - $y = \text{factor} * v$



8

Example Mapping - Scale


- Rotate by θ degrees
 - $x = u \cdot \cos(\theta) - v \cdot \sin(\theta)$
 - $y = u \cdot \sin(\theta) + v \cdot \cos(\theta)$




9

Example Mapping - Shear

- Shear in X by *factor*
 - $x = u + \text{factor} \cdot v$
 - $y = v$



- Shear in Y by *factor*
 - $x = u$
 - $y = v + \text{factor} \cdot u$



10

Other Mappings

- Any function of u and v :
 - $x = f_x(u,v)$
 - $y = f_y(u,v)$




11

Implementing Warping - 1


- Forward Mapping

```
for (int u = 0; u < umax; u++) {
  for (int v = 0; v < vmax; v++) {
    float x = fx(u,v);
    float y = fy(u,v);
    dest(x,y) = src(u,v);}}
```



Forward Mapping – Bad Idea!!

- Iterating over source image creates issues
 - Double coverage & Empty pixels




13

Better Warping

- Reverse Mapping


```
for (int x = 0; x < xmax; x++) {
  for (int y = 0; y < ymax; y++) {
    float u =  $f_x^{-1}(x,y)$ ;
    float v =  $f_y^{-1}(x,y)$ ;
    dest(x,y) = src(u,v);}
```



14

Reverse Mapping


- Iterate over destination image
- May oversample, but much better!



15

Reverse Mapping (cont.)

- Evaluate arbitrary, *non-integer* (u,v) coords
- Resampling!!!



16

Outline

- Warping
 - Scale, Rotate, Warps
- **Sampling & Reconstruction**
 - **Resampling**
 - **Reconstruction Filters**
- Image Manipulation
 - Filtering
 - Pixel Operation
- Quantization

17

Signal Resampling

- Changing the sampling rate
 - Images: Enlarging or shrinking!
- Creating more samples:
 - Increasing sampling rate
 - “Upsampling”
 - “Enlarging”
- Ending up with fewer samples:
 - Decrease the sampling rate
 - “Downsampling”
 - “Reducing”

18


Reducing and Enlarging

- Very common operation
 - Different device resolution (mobile vs HDTV)
 - Different memory availability
- Very commonly done poorly
- Simple approach: Drop/replicate pixels
- Correct approach: Use resampling
 - Step 1: Reconstruct image
 - Step 2: Sample reconstructed image

19

Signal Reconstruction

- Continuous-discrete convolution



20

Reconstruction Pseudocode

- Reconstruction is just continuous discrete filtering

```

function reconstruct(sequence  $a$ , filter  $f$ , real  $x$ )
   $s = 0$ 
   $r = f.\text{radius}$ 
  for  $i = \lceil x - r \rceil$  to  $\lfloor x + r \rfloor$  do
     $s = s + a[i]f(x - i)$ 
  return  $s$ 


```

- Resamples the signal at point x
- Different filters give different results

21

Resampling

- Reconstruction creates a continuous function
 - Sample *this* new function




22

Cont.-Disc. Convolution in 2D

- Same convolution, but with two variables

$$(a \star f)(x, y) = \sum_{i,j} a[i, j] f(x - i, y - j)$$



- Loop over nearby pixels
 - Compute weights with $f(x, y)$; Take weighted avg.
- Looks like discrete filter, but:
 - Offsets are not integers
 - Filter is continuous
- Remember: Filter will shift relative to the grid
 - Can evaluate at any point



23

Cont.-Disc. Convolution in 2D

$$(a \star f)(x, y) = \sum_{i,j} a[i, j] f(x - i, y - j)$$



24

Reconstruction Filters

(A Gallery)

- Box filter
 - Simple & cheap
- Tent filter
 - Linear interpolation
- Gaussian filter
 - Very smooth antialiasing filter
- B-spline cubic
 - Very smooth
- Catmull-rom cubic
 - Interpolating
- Mitchell-Netravali cubic
 - Good for upsampling


25

Box Filter

- Simple average
 - Uniform weights

$$a_{\text{box},r}[i] = \begin{cases} 1/(2r+1) & |i| \leq r, \\ 0 & \text{otherwise.} \end{cases}$$

$$f_{\text{box},r}(x) = \begin{cases} 1/(2r) & -r \leq x < r, \\ 0 & \text{otherwise.} \end{cases}$$




26

Tent Filter

- Linear interpolation
 - Distance dependent weights

$$f_{\text{tent}}(x) = \begin{cases} 1 - |x| & |x| < 1, \\ 0 & \text{otherwise;} \end{cases}$$


$$f_{\text{tent},r}(x) = \frac{f_{\text{tent}}(x/r)}{r}.$$



27

Gaussian Filter

- Smoothed interpolation
 - Distance dependent weights




$$f_g(x) = \frac{1}{\sqrt{2\pi}} e^{-x^2/2}.$$

28

B-Spline Cubic

- Smoothed interpolation
 - Distance dependent weights
 - Simpler to compute the Gaussian




$$f_B(x) = \frac{1}{6} \begin{cases} -3(1 - |x|)^3 + 3(1 - |x|)^2 + 3(1 - |x|) + 1 & -1 \leq x \leq 1, \\ (2 - |x|)^3 & 1 \leq |x| \leq 2, \\ 0 & \text{otherwise.} \end{cases}$$

29

Catmull-Rom

- Has negative values!
 - Provides some sharpening to counteract blur





$$f_C(x) = \frac{1}{2} \begin{cases} -3(1 - |x|)^3 + 4(1 - |x|)^2 + (1 - |x|) & -1 \leq x \leq 1, \\ (2 - |x|)^3 - (2 - |x|)^2 & 1 \leq |x| \leq 2, \\ 0 & \text{otherwise.} \end{cases}$$

30

Michell-Netravali Cubic

- Created a parameterized cubic curve
 - Held a user study to find best parameters



$$f_M(x) = \frac{1}{3}f_B(x) + \frac{2}{3}f_C(x)$$

$$= \frac{1}{18} \begin{cases} -21(1-|x|)^3 + 27(1-|x|)^2 + 9(1-|x|) + 1 & -1 \leq x \leq 1, \\ 7(2-|x|)^3 - 6(2-|x|)^2 & 1 \leq |x| \leq 2, \\ 0 & \text{otherwise.} \end{cases}$$

31

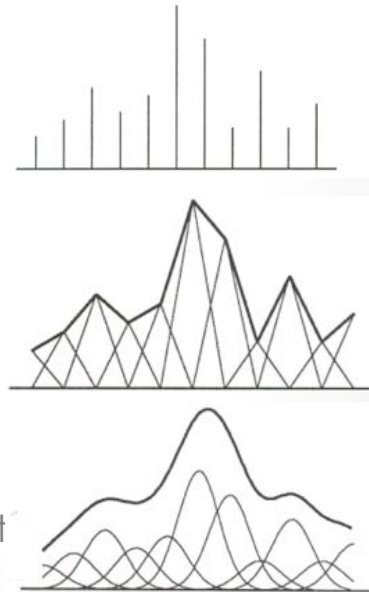

Hacks as Filters

- Some simple approaches are same as filters
- Box filter (radius 0.5): Nearest Neighbors
 - Will only catch the one closest input
 - $\text{output}[i, j] = \text{input}[\text{round}(x(i)), \text{round}(y(j))]$
- Tent filter (radius 1): Linear Interpolation
 - Will catch exactly 2 points
 - Weights are a and $(1-a)$
 - Results in straight line interpolation of points

32

Common in Practice

- Box – Cause Jaggies
- Bilinear filtering (Tent)





- Gaussian filtering
 - Width of Gaussian affects blurring amount

33

Ring, Overshoot, & Ripples

- Overshoot
 - Caused by negative filter values
- Ringing
 - Reconstructing a straight line may not be constant







34

$$\sum_i f(x+i) = 1 \quad \text{for all } x.$$


Making 2D Filters

- Compute dimensions independently (separable filter)



35

Upsampling Illustrated



36

Implementing Scale

- Scale (src, dst, sx, sy)

```
float r = max(1.0/sx, 1.0/sy);
for (int x = 0; x < xmax; x++) {
    for (int y = 0; y < ymax; y++) {
        float u = x / sx;
        float v = y / sy;
        dest(x,y) = resample_src(u,v,r);
    }
}
```

37

Implementing Rotate

- Rotate(src, dst, theta)

```
for (int x = 0; x < xmax; x++) {
    for (int y = 0; y < ymax; y++) {
        x = u*cos(-theta) - v*sin(-theta)
        y = u*sin(-theta) + v*cos(-theta)
        dest(x,y) = resample_src(u,v,r);
    }
}
```



38

Resampling Example



1000 pixel width

[Philip Greenspun]



by dropping pixels

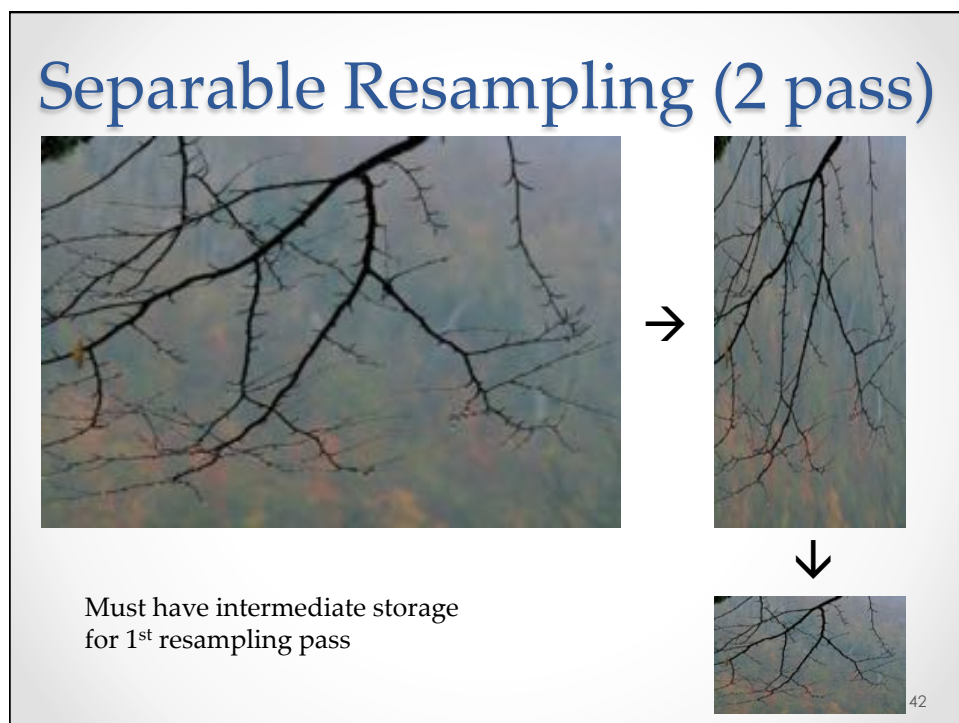



gaussian filter

250 pixel width

[Philip Greenspun]

40



Artifacts

- Jaggies, Moiré patterns, etc.


43



[Hearn & Baker cover]

600ppi scan of a color halftone image

44



Outline

- Warping
 - Scale, Rotate, Warps
- Sampling & Reconstruction
 - Resampling
 - Reconstruction Filters
- **Image Manipulation**
 - **Filtering**
 - Pixel Operation
- Quantization

Filter Examples


- Blur (and Sharpen)



Original



Blur



Sharpen

Filter =

$$\begin{bmatrix} 1/16 & 2/16 & 1/16 \\ 2/16 & 4/16 & 2/16 \\ 1/16 & 2/16 & 1/16 \end{bmatrix}$$

Edge Detect



Original




Detected Edges

$$\text{Filter} = \begin{bmatrix} -1 & -1 & -1 \\ -1 & +8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$$

Ugly Details

- What about near the edge?
 - The filter window falls off the edge
 - Must extrapolate
 - Methods:
 - clip filter (black)
 - wrap around
 - copy edge
 - reflect across edge
 - vary filter near edge



[Philip Greenspun]

49

Outline

- Warping
 - Scale, Rotate, Warps
- Sampling & Reconstruction
 - Resampling
 - Reconstruction Filters
- **Image Manipulation**
 - Filtering
 - **Pixel Operation**
- Quantization

50

Brighten


- Simply scale pixel value
 - Must clamp to range (e.g., 0 to 255)
- Trick: Interpolate/extrapolate from all black img.



Original



Brighten




51

Contrast


- Compute mean luminance L over all pixels
 - Luminance = $0.30*r + 0.59*g + 0.11*b$
- Scale deviation from L for each pixel
- Trick: Interpolate/extrapolate from avg. grey img.



Original



More Contrast



52

Saturation

- Interpolate/extrapolate from grayscale version of image



53


Outline

- Warping
 - Scale, Rotate, Warps
- Sampling & Reconstruction
 - Resampling
 - Reconstruction Filters
- Image Manipulation
 - Filtering
 - Pixel Operation
- **Quantization**

54

Quantization




- Artifacts due to limited intensity resolution
 - Frame buffers have limited bits/pixel
 - Physical devices have limited dynamic range



55

Uniform Quantization

- $P(x,y) = \text{trunc}(I(x,y) + .5)$


 $I(x,y)$

 $P(x,y)$
(2 bits per pixel)

56

Uniform Quantization

- Effect of decreasing bits/pixel



8 bits

4 bits

2 bits

1 bit

57


Reducing Quantization Effects

- Halftoning
 - Classical Halftoning
- Dithering
 - Random Dithering
 - Ordered Dithering
 - Error Diffusion Dithering

58

Classical Halftoning


- Use dots of varying size to represent intensity
 - Area of dot proportional to intensity



59


Halftone Patterns

- Use clusters of pixels to represent intensity
 - Trades spatial resolution for intensity resolution



60

Halftoning Examples



61

Dithering


- Distribute errors among pixels
 - Exploit eye's natural spatial integration
 - Display larger range of perceptible intensities



Original
(8 bits)



Uniform
Quantization
(1 bit)



Floyd-Steinberg
Dithering
(1 bit)

62

Random Dithering

- Randomize quantization errors
 - Errors appear as noise

$$P(x,y) = \text{trunc}(I(x,y) + \text{noise}(x,y) + .5)$$



Original
(8 bits)



Uniform Quantization
(1 bit)



Random Dither
(1 bit)

63

Ordered Dither

- Structured quantization error
 - Matrix store pattern of thresholds

```

i = x mod n
j = y mod n
e = I(x,y) - trunc(I(x,y))
if (e > D(i,j))
    P(x,y) = ceil(I(x, y))
else
    P(x,y) = floor(I(x,y))
  
```

$$D_2 = \begin{bmatrix} 3 & 1 \\ 0 & 2 \end{bmatrix}$$

64

Ordered Dither

- Bayer's ordered dither matrices

$$D_n = \begin{bmatrix} 4D_{n/2} + D_2(1,1)U_{n/2} & 4D_{n/2} + D_2(1,2)U_{n/2} \\ 4D_{n/2} + D_2(2,1)U_{n/2} & 4D_{n/2} + D_2(2,2)U_{n/2} \end{bmatrix}$$

$$D_2 = \begin{bmatrix} 3 & 1 \\ 0 & 2 \end{bmatrix} \quad D_4 = \begin{bmatrix} 15 & 7 & 13 & 5 \\ 3 & 11 & 1 & 9 \\ 12 & 4 & 14 & 6 \\ 0 & 8 & 2 & 10 \end{bmatrix}$$

65

Ordered Dither - Compare



Original
(8 bits)


Uniform Quantization
(1 bit)

Ordered Dither
(1 bit)

66


Error Diffusion Dither

- Spread quantization error over neighboring pixels
 - Error dispersed to pixels right and below



67

Dither Comparison



68

Image Operations

- Shrink/Enlarge
- Blur/Sharpen
- Edge detection
- Brighten/Darken
- Change Contrast
- Saturate/Desaturate
- Quantize
- Random Dither
- Ordered Dither
- Floyd-Stienberg

69

Assignment 1

- Shrink/Enlarge
- Blur/Sharpen
- Edge detection
- Brighten/Darken
- Change Contrast
- Saturate/Desaturate
- Quantize
- Random Dither
- Ordered Dither
- Floyd-Stienberg

70

Assign. 1 - Details

- Lots of code given already:
 - Read/Write a images (stb_image_write.h)
 - Command line processing
 - Brighten
 - Image & Pixel class
 - Bayer constants
 - Floyd-Stienberg constants
- Should compile very easily!
- Due two weeks from now.

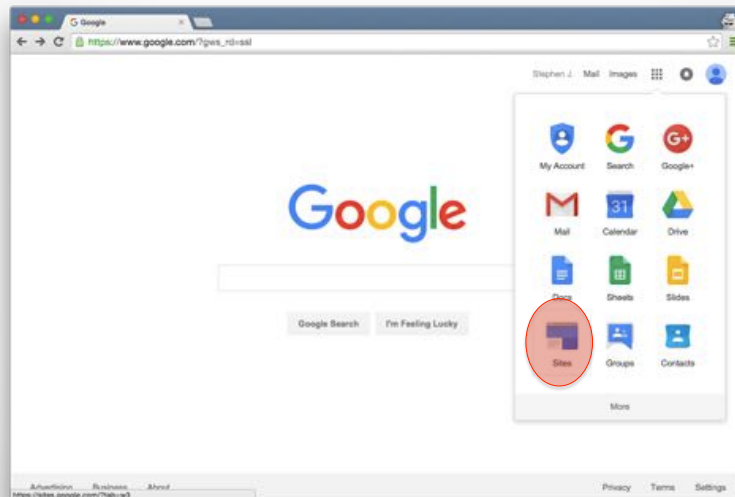
71

HW 0 Turn-in

- What to put on webpage:
 - Code & Any libraries needed to compile
 - Short write-up
 - Pictures of working project!
- Turn in webpage:
 - Through Moodle link
- Create Webpage
 - Easy: Just place html on a server
 - Very Easy: Google Sites through UMN

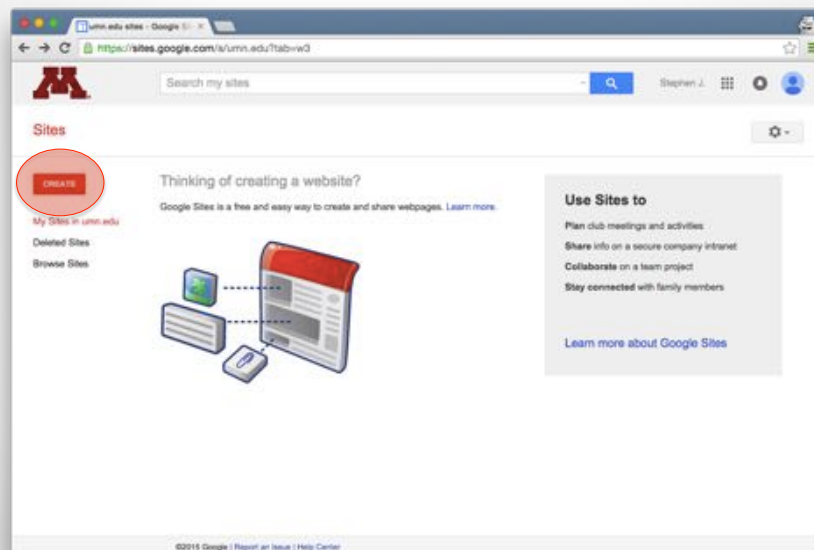
72

Creating a Google Site - 1



73

Creating a Google Site - 2



74

Creating a Google Site - 3



75

Permission

- If you use Google Sites (or a similar service):
 - Make sure you share it university-wide
- Ensures both me and the TA's can access the webpage

76